










Reading the diagram:

- **AI sits at the bottom**, as a constrained component.
- **Hardware-rooted trust + manifest** define what is allowed.
- **Zero-trust + enforcement** ensure rules cannot be changed at runtime.
- **Scenario & threat modelling** shapes the safety architecture.
- **Safety & ethics** form the hard boundary for all behaviour.
- **Applications** interact only through this layered safety stack.